



CIS 263 Introduction to Data Structures and Algorithms

B-Trees

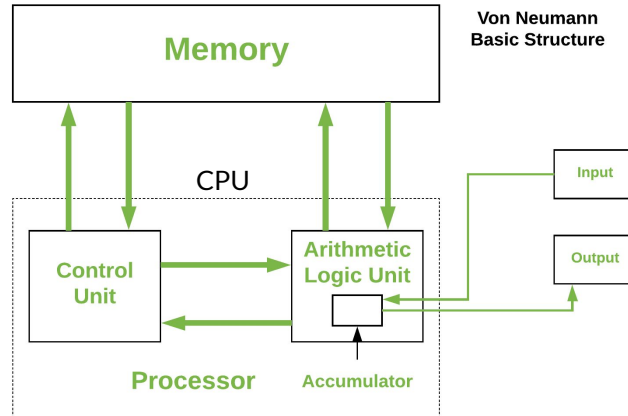


B-Tree

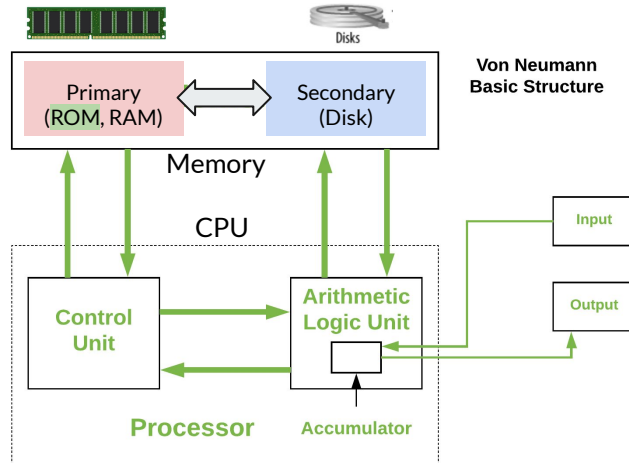
- A **B-tree** is a self-balancing tree.
- The B-tree generalizes BST, allowing more than two children.
- Unlike other self balancing (AVL Tree, RBT), the B-Tree is well suited for **storage systems** that read and write relatively large blocks of data, such as [databases](#) and [file systems](#).

<https://en.wikipedia.org/wiki/B-tree>

Computer Architecture basics

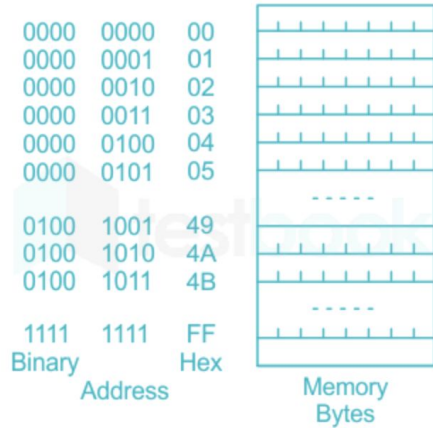


Computer Architecture basics



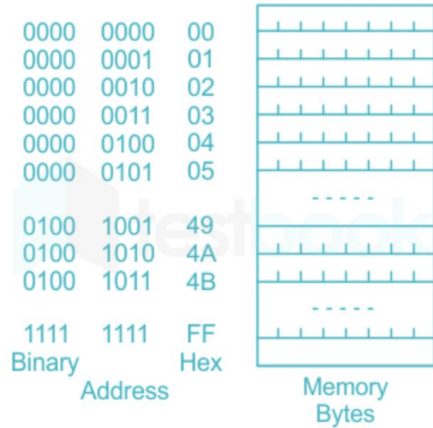
- In this course, our main focus will be mainly focused on the **memory management** part; more importantly the **primary memory**
- And **number of disk accesses** (latency delay)
- Through
 - Using appropriate data structures, and
 - Algorithms
- We are not doing any improvements over architecture etc.

Memory, the concept of Array



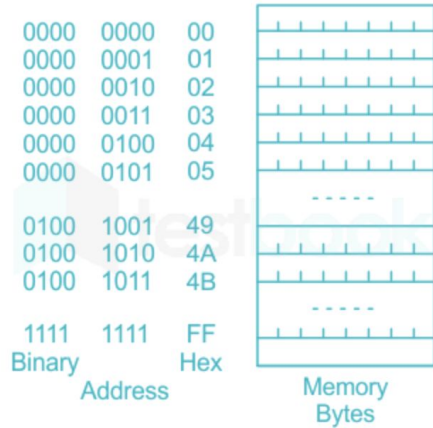
- 8 bit memory (a simple case)

Memory, the concept of Array



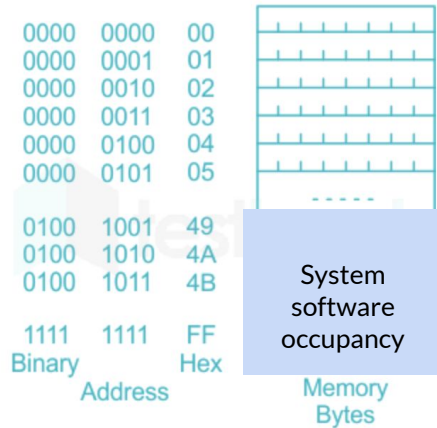
- 8 bit memory (a simple case)
- What's the total memory size?

Memory, the concept of Array



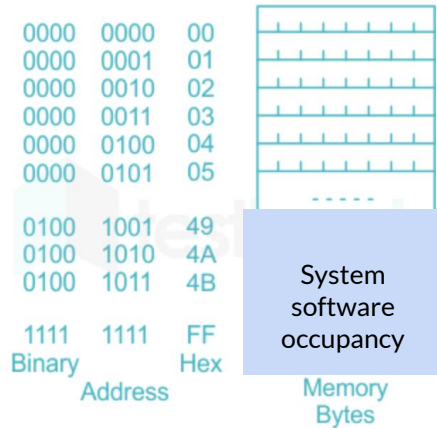
- 8 bit memory (a simple case)
- What's the total memory size?
 - 2^8 or 1024 byte

Memory, the concept of Array



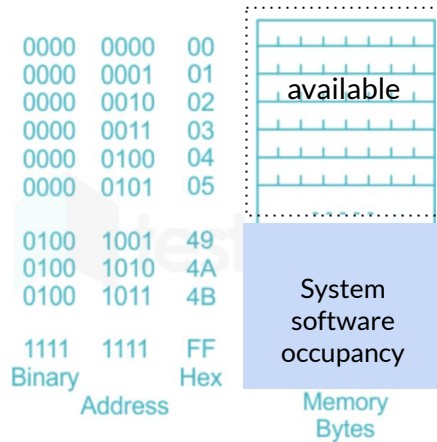
- 8 bit memory (a simple case)
- What's the total memory size?
 - 2^8 or 1024 byte
- Addressing is inherent (consecutive, a sequence)

Memory, the concept of Array



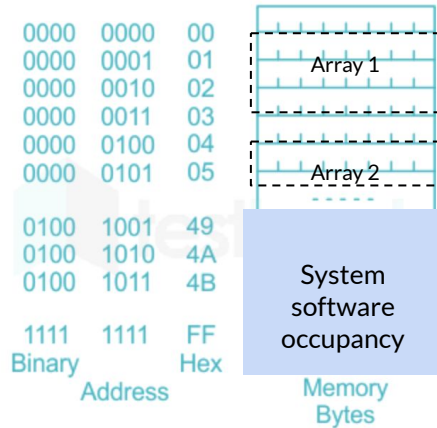
- 8 bit memory (a simple case)
- What's the total memory size?
 - 2^8 or 1024 byte
- Addressing is inherent (consecutive, a sequence)
- Let's assume our system softwares reserved part of the memory like

Memory, the concept of Array



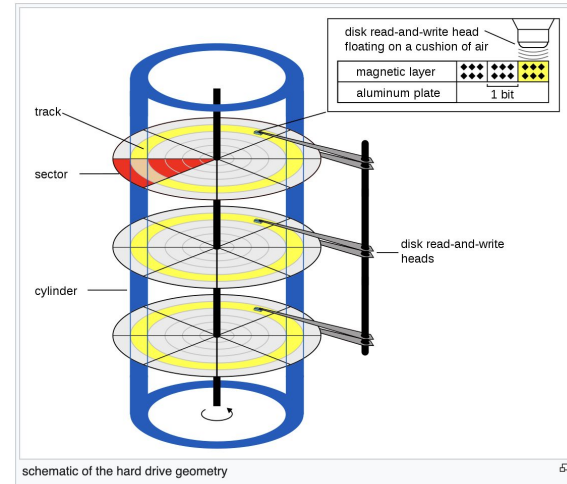
- Addressing is inherent (consecutive, a sequence)
- Let's assume our system softwares reserved part of the memory like

Memory, the concept of Array

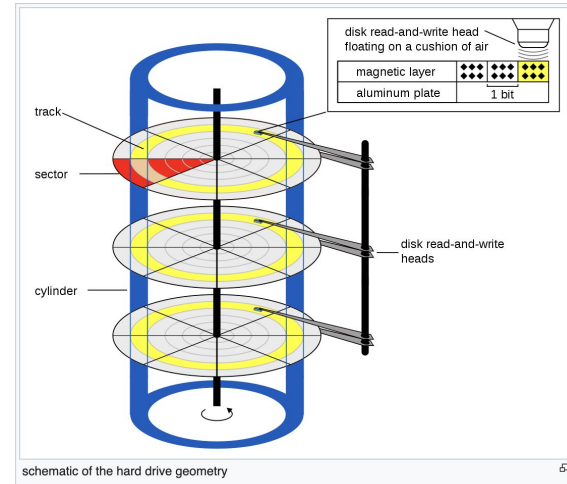


- Addressing is inherent (consecutive, a sequence)
- Let's assume our system softwares reserved part of the memory like
- **Array** (our first data structure) meaning holding a fixed block whether you use it or not.

Disk: track, sector, block



Disk: track, sector, block, offset





It's all about how to reduce number of disk accesses



Disk: track, sector, block, offset

An exercise:

- One student record requires 128 Byte
- **How many such records a block can store at max ?**

One student record requires 128 Byte

Id	Name	Dept	Class
0			
1			
...			
n-2			
n-1			



Disk: track, sector, block, offset

An exercise:

- One student record requires 128 Byte
- How many such records a block can store at max ?
- $512 / 128 = 4$

One student record requires 128 Byte

Id	Name	Dept	Class
0			
1			
...			
n-2			
n-1			



Disk: track, sector, block, offset

An exercise:

- One student record requires 128 Byte
- How many such records a block can store at max ?
- $512 / 128 = 4$
- How many blocks you will require to store $n=100$ student records?

One student record requires 128 Byte

Id	Name	Dept	Class
0			
1			
...			
n-2			
n-1			



Disk: track, sector, block, offset

An exercise:

- One student record requires 128 Byte
- How many such records a block can store at max ?
- $512 / 128 = 4$
- How many blocks you will require to store $n=100$ student records?
- $100 / 4 = 25$ blocks; right?

One student record requires 128 Byte

Id	Name	Dept	Class
0			
1			
...			
n-2			
n-1			

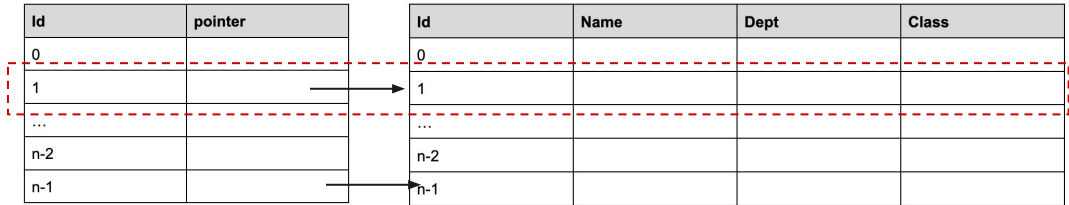
How about adding more content

An exercise:

- One student record requires 128 Byte
- 16 additional bytes for the pointer record
- **How many pointer records a Block can store?**

One student record requires 128 Bytes

One (id + pointer takes): 16 Bytes



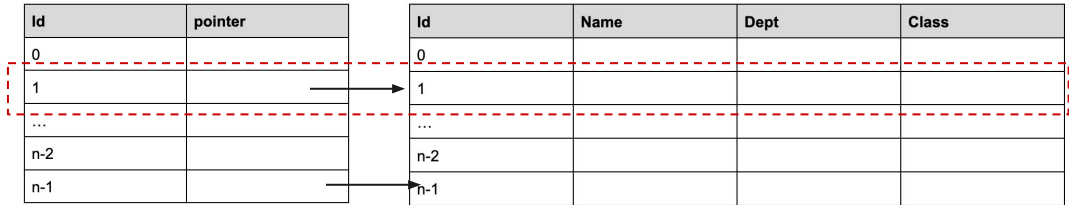
How about adding more content

An exercise:

- One student record requires 128 Byte
- 16 additional bytes for the pointer record
- How many pointer records a Block can store?
- $512 / 16 = 32$

One student record requires 128 Bytes

One (id + pointer takes): 16 Bytes



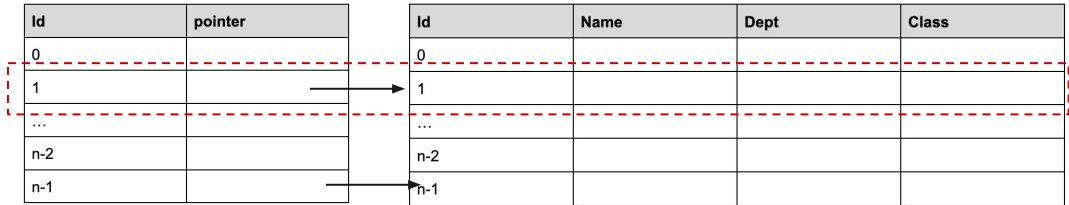
How about adding more content

An exercise:

- One student record requires 128 Byte
- 16 additional bytes for the pointer record
- How many pointer records a Block can store?
- $512 / 16 = 32$
- How many block you will require to query 100 records ?

One student record requires 128 Bytes

One (id + pointer takes): 16 Bytes



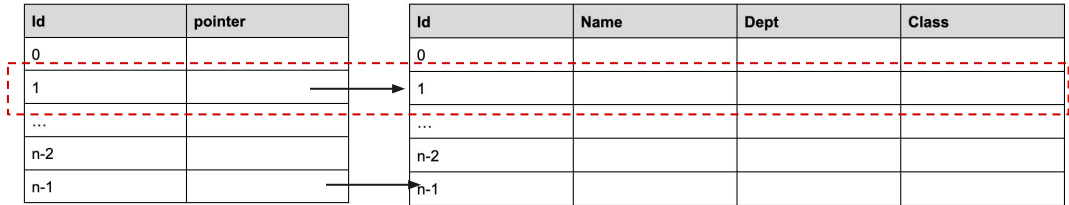
How about adding more content

An exercise:

- One student record requires 128 Byte
- 16 additional bytes for the pointer record
- How many pointer records a Block can store?
- $512 / 16 = 32$
- How many block you will require to query 100 records ?
- $100 / 32 = \sim 4$ (max)

One student record requires 128 Bytes

One (id + pointer takes): 16 Bytes



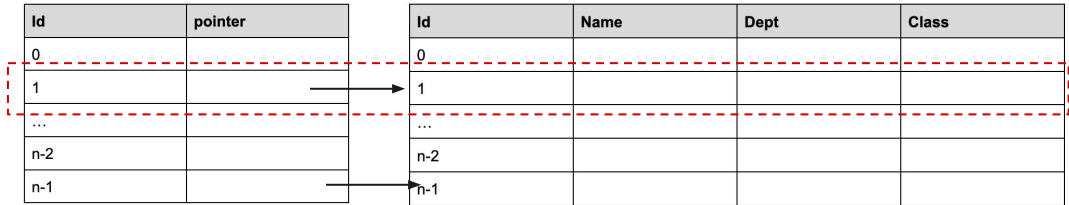
How about adding more content

An exercise:

- One student record requires 128 Byte
- 16 additional bytes for the pointer record
- How many pointer records a Block can store?
- $512 / 16 = 32$
- How many block you will require to query 100 records ?
- $100 / 32 = \sim 4$ (max)
- **So we reduced from 25 - > 4 Block accesses; right?**
- **This is the general idea of database indexing.**

One student record requires 128 Bytes

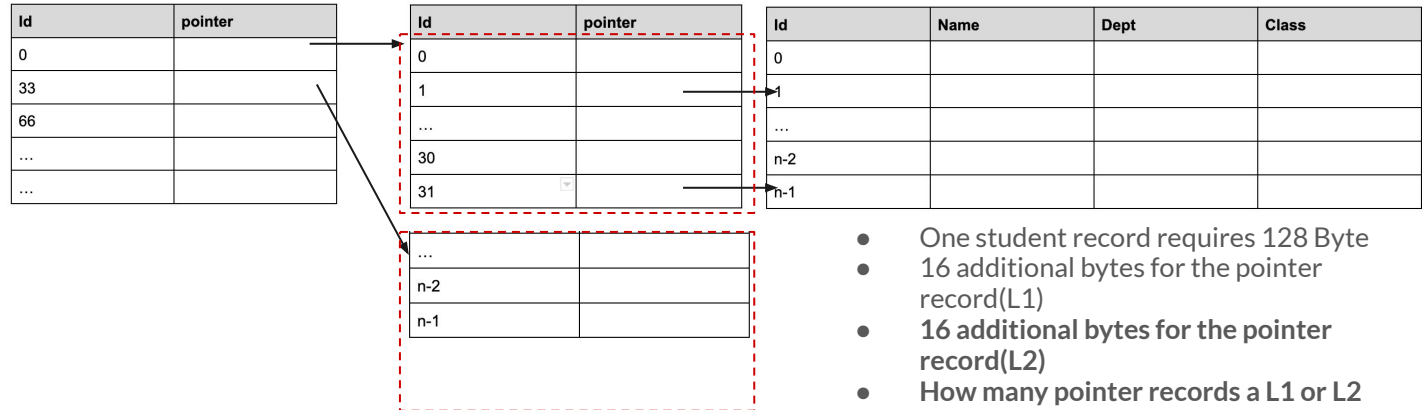
One (id + pointer takes): 16 Bytes



How about adding more content

One student record requires 128 Bytes

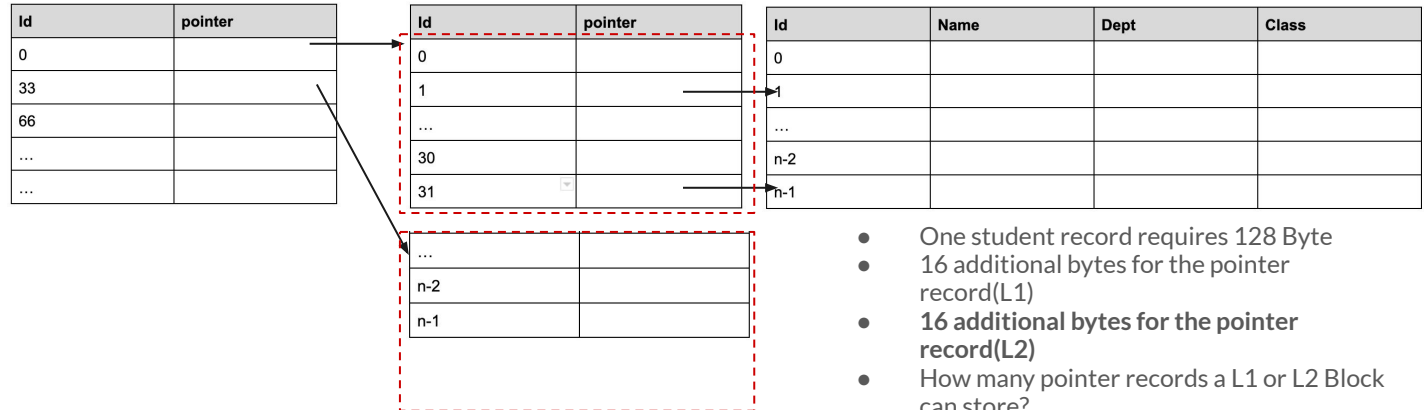
One (id + pointer takes): 16 Bytes



How about adding more content

One student record requires 128 Bytes

One (id + pointer takes): 16 Bytes

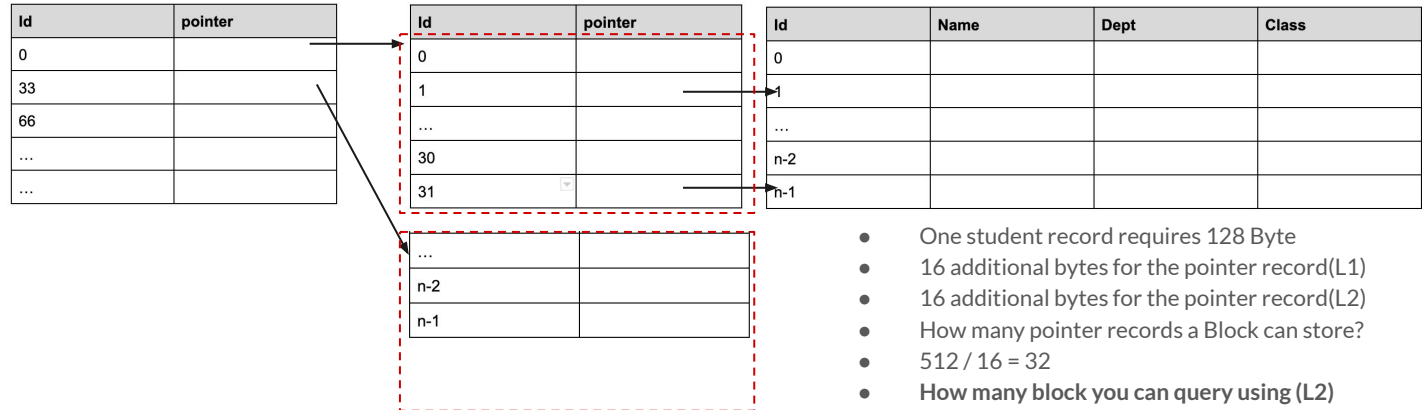


- One student record requires 128 Byte
- 16 additional bytes for the pointer record(L1)
- **16 additional bytes for the pointer record(L2)**
- How many pointer records a L1 or L2 Block can store?
- **512**

How about adding more content

One student record requires 128 Bytes

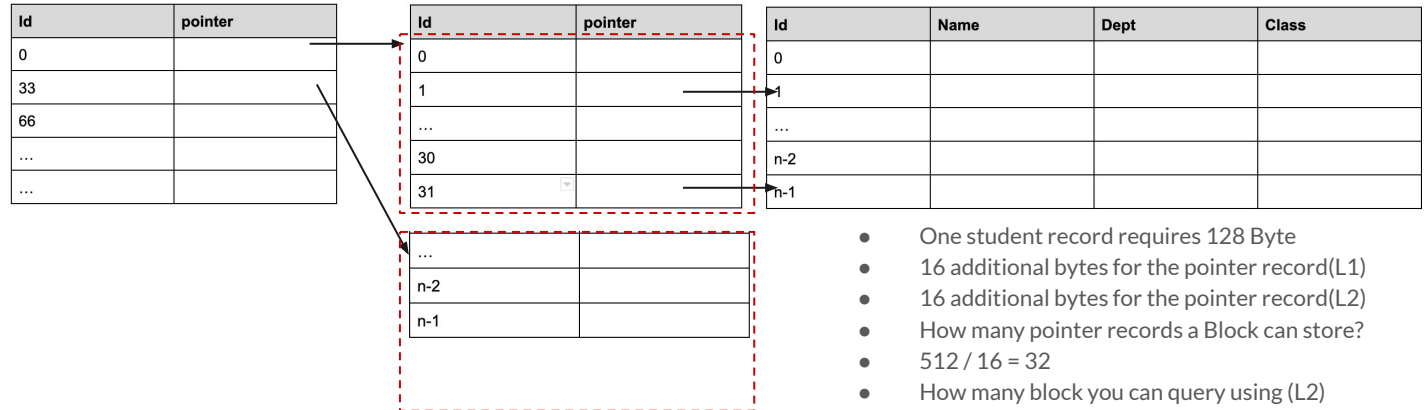
One (id + pointer takes): 16 Bytes



How about adding more content

One student record requires 128 Bytes

One (id + pointer takes): 16 Bytes





m-way tree

- $m-1$ keys
- m children



B-tree

- Every node has $m/2$ (*ceil*) children at least
- Root has minimum **2 children**
- All leaves at same level
- Bottom up method



Generating a m(3)-way B-Tree

- $M = 3$
- $\text{Max } 3 - 1 = 2$ keys per node

45

Input sequence: 45



Generating a m(3)-way B-Tree

- $M = 3$
- $\text{Max } 3 - 1 = 2$ keys per node

35, 45

Input sequence: 45, 35



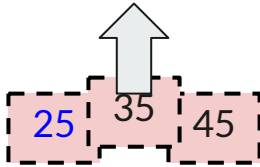
Generating a m(3)-way B-Tree

25, 35, 45

- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds

Input sequence: 45, 35, 25

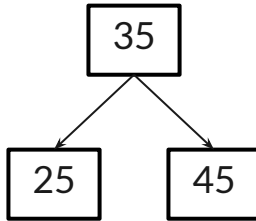
Generating a m(3)-way B-Tree



- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds
- Split and adjust levels

Input sequence: 45, 35, 25

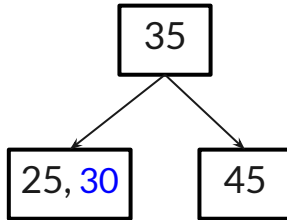
Generating a m(3)-way B-Tree



- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds
- Split and adjust levels

Input sequence: 45, 35, 25

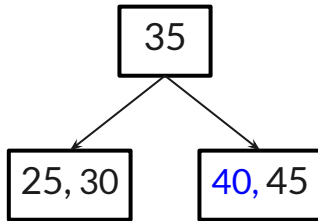
Generating a m(3)-way B-Tree



- $M = 3$
- $\text{Max } 3-1 = 2$ keys per node

Input sequence: 45, 35, 25, 30

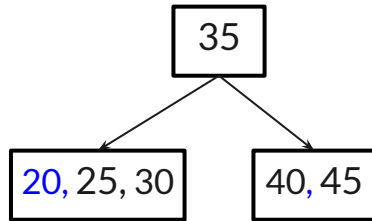
Generating a m(3)-way B-Tree



- $M = 3$
- $\text{Max } 3 - 1 = 2$ keys per node

Input sequence: 45, 35, 25, 30, 40

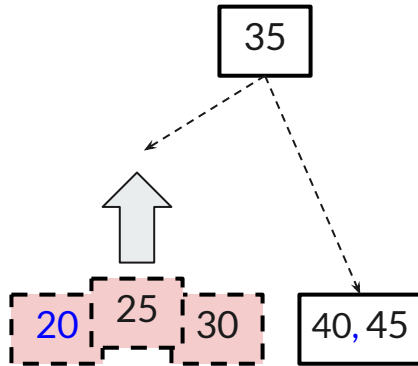
Generating a m(3)-way B-Tree



- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds
- Split and adjust levels

Input sequence: 45, 35, 25, 30, 40, 20

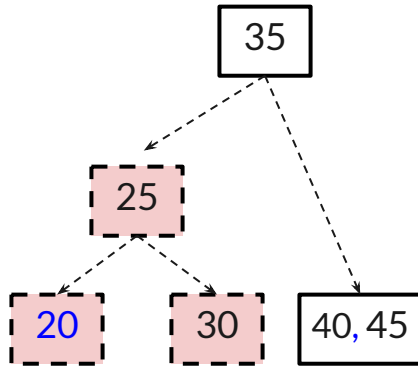
Generating a m(3)-way B-Tree



- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds
- Split and adjust levels

Input sequence: 45, 35, 25, 30, 40, 20

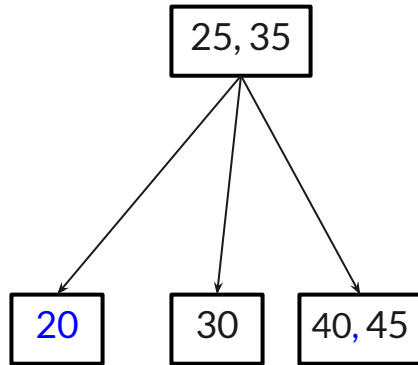
Generating a m(3)-way B-Tree



- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds
- Split and adjust levels

Input sequence: 45, 35, 25, 30, 40, 20

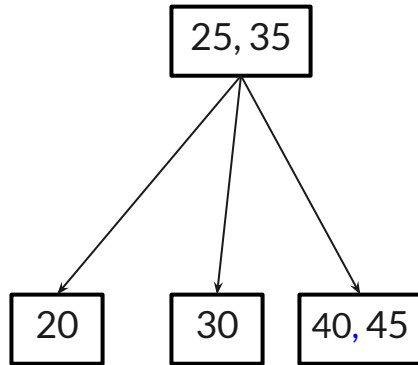
Generating a m(3)-way B-Tree



- $M = 3$
- Max $3 - 1 = 2$ keys per node
- Limit exceeds
- Split and adjust levels

Input sequence: 45, 35, 25, 30, 40, 20

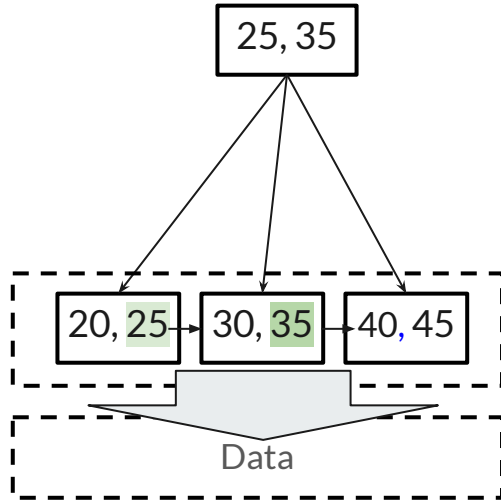
Generating a m(3)-way B-Tree



- $M = 3$
- $\text{Max } 3-1 = 2 \text{ keys per node}$

Input sequence: 45, 35, 25, 30, 40, 20

B+ Tree



- $M = 3$
- $\text{Max } 3 - 1 = 2$ keys per node
- All keys (both internal and leaf nodes) are available at the leaf level (duplicate copies of root/intermediate nodes at the leaf level)
- Leaf nodes contain data pointers